

IMAGE GENERATION WITH DIFFUSION MODELS

Attila Nemes (B6RYIK),
Csaba Potyok (OZNVQ4)



Agenda

- review of literature (tutorial video watching)
- acquisition and preparation of datasets (Flowers102, CelebA)
- baseline model (VAE)
- designing and training a DDPM model
- evaluation of the results
- creating an AI service

Datasets



Flowers102 (~8k images)




CelebA (~200k images)

(source: <https://huggingface.co/datasets/nielsr/CelebA-faces>)

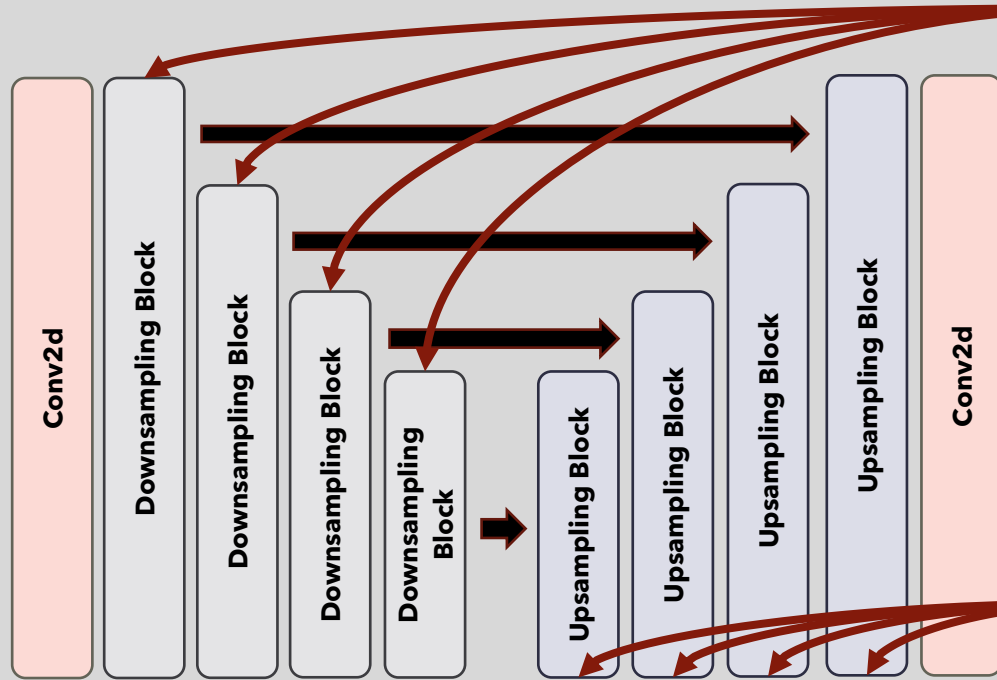
2024. 12. 12.

Baseline model

- simple VAE
- 10 epoch training
- no hyperparameter optimization




DDPM architecture

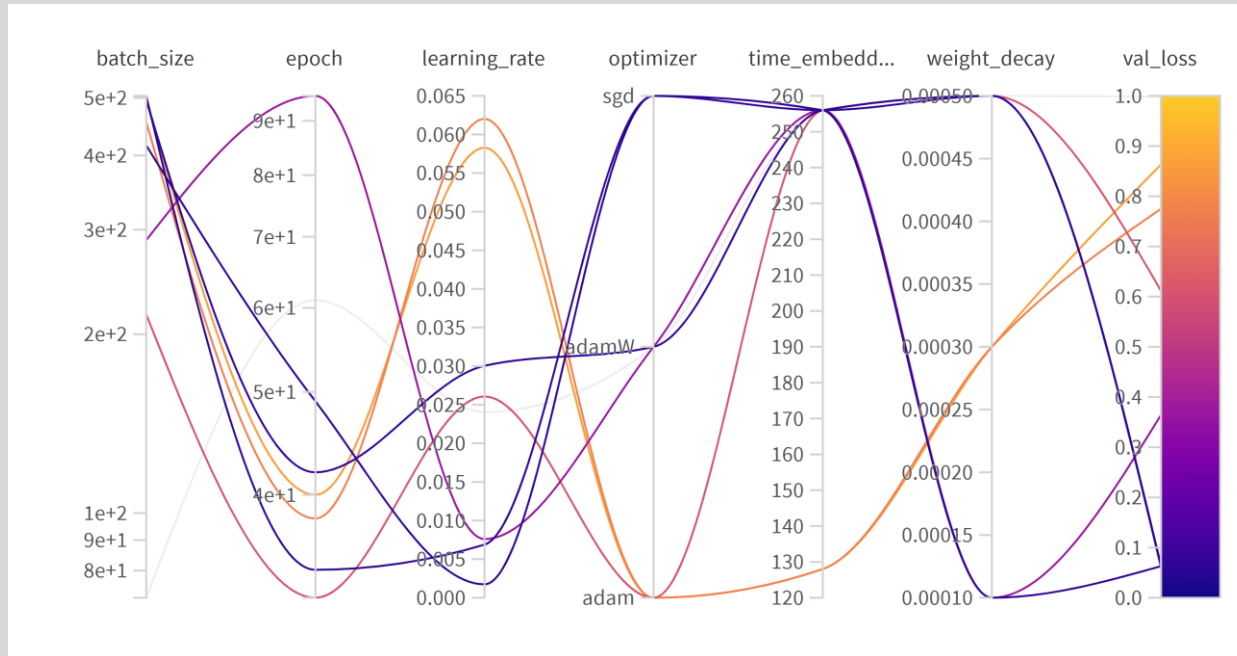


Training models

- using Pytorch + Lightning
- LightningModule for pipeline
- Diffusion model
- running on Komondor



Hyperparameter optimization



- using wandb
- Bayesian hyperopt
- for both datasets
- running on Komondor

Evaluation of models

model	FID	IS mean	IS std	loss
1	60,68	1,328	0,021	0,057
2	19,33	1,527	0,026	0,059
3	100,41	2,485	0,123	0,055
4	17,958	1,579	0,054	0,062
5	14,353	1,738	0,063	0,065
6	15,57	1,562	0,029	0,077
7	8,477	1,628	0,048	0,067
8	25,804	1,7	0,048	0,073

Result on Flowers102

model	FID	IS mean	IS std	loss
1	26,072	1,568	0,006	0,04
2	18,794	1,186	0,003	0,037
3	16,499	1,605	0,008	0,041
4	16,39	1,374	0,0048	0,039
5	16,606	1,349	0,003	0,038

Result on CelebA

Generated images



Best with Flowers102



Best with CelebA

Demo video

